

Introduction to Shading

Contents

Multiple Lights

Linear Light Response

Imagine that you want to photograph an object such as the toy in figure 1 with different lights. Let's say that all lights are turned off to start with. You then switch on the first light, take a picture, switch the light off and repeat this process for each light. At the end, turn all the lights on and take another picture. If you now add up all the pictures that you took with each light turned on individually and compare it to the picture where all the lights were turned on at once, then the two resulting images should look the same.

$$\text{Image}_{\text{light1}} + \text{Image}_{\text{light2}} + \text{Image}_{\text{light3}} + \dots = \text{Image}_{\text{all lights}}.$$

In other words, the contribution of each light adds up linearly. This is an important observation for two reasons. First if you wish to create photorealistic images then it is important for your renderer to follow the same principle. It is also important for artists recomposing CG renders from individual layers where each layer represents the contribution of one particular light from the scene. Adding a curve to change the brightness of the image in a non-linear way (such as adding a gamma for instance) before adding it up to the other layers would lead to an incorrect result from a physical point of view. More importantly from a programming point of view, this principle simply means that in the renderer, the contribution of each light just needs to be summed up. In other words, the total amount of light arriving on a point is just the linear sum of the amount of light that each light is contributing to.



Figure 1: the contribution of each light adds up linearly.

In mathematical terms, for a diffuse surface, this concept can be written using the following formula:

$$S_P = \sum_{N=0}^{N=(\text{nlights}-1)} \frac{\rho_d}{\pi} * Li_N * \cos(N \cdot L_N).$$

Where S_P stands for shading point. The symbol \sum in mathematics means "sum". In other words, for each light in the scene (there is *nlights* in total), we compute the diffuse equation by replacing in the equation the term Li_N and L_N by the current light intensity and direction. This is the same as:

$$S_P = \frac{\rho_d}{\pi} * (Li_0 * \cos(N \cdot L_0) + Li_1 * \cos(N \cdot L_1) + \dots + Li_{N-1} * \cos(N \cdot L_{N-1})).$$

From a programming point of view, handling more than one light source is very simple. First you can store all the lights in a list and pass this list to the `castRay()` function which is where shading is done. We then iterate over all the lights and add their contribution to the shaded point illumination. Keep in mind that the light contribution is attenuated by the cosine of the angle between P 's normal and the light direction. This term is different for each light and thus needs to be computed for each light. Similarly a shadow ray needs to be cast for each light as well. In the end, the code looks as follows:

```
001 int main(int argc, char **argv)
002 {
003     ...
004     // loading geometry
005     std::vector<std::unique_ptr<Light>> lights;
```

```

006
007     Matrix44f l2w1;
008     l2w1[3][0] = 2;
009     l2w1[3][1] = 4;
010
011     lights.push_back(std::unique_ptr<Light>(new PointLight(l2w1, Vec3f(1, 0.6, 0.6), 500)));
012
013     Matrix44f l2w2;
014     l2w2[3][0] = -1;
015     l2w2[3][1] = 4;
016     l2w2[3][2] = -1;
017
018     lights.push_back(std::unique_ptr<Light>(new PointLight(l2w2, Vec3f(0.6, 0.6, 1), 500)));
019
020     // finally, render
021     render(options, objects, lights);
022     ...
023 }
024
025 Vec3f castRay(
026     const Vec3f &orig, const Vec3f &dir,
027     const std::vector<std::unique_ptr<Object>> &objects,
028     const std::vector<std::unique_ptr<Light>> &lights,
029     const Options &options)
030 {
031     Vec3f hitColor = 0;
032     IsectInfo isect;
033     if (trace(orig, dir, objects, isect)) {
034         Vec3f hitPoint = orig + dir * isect.tNear;
035         Vec3f hitNormal;
036         Vec2f hitTexCoordinates;
037         isect.hitObject->getSurfaceProperties(hitPoint, dir, isect.index,
038         isect.uv, hitNormal, hitTexCoordinates);
039         for (uint32_t i = 0; i < lights.size(); ++i) {
040             Vec3f lightDir, lightIntensity;
041             IsectInfo isectShad;
042             lights[i]->getShadingInfo(hitPoint, lightDir, lightIntensity, isectShad.tNear);
043             bool vis = !trace(hitPoint + hitNormal * options.bias, -lightDir,
044             objects, isectShad, kShadowRay);
045             // accumulate light contribution
046             hitColor += vis * isect.hitObject->albedo * lightIntensity *
047             std::max(0.f, hitNormal.dotProduct(-lightDir));
048         }
049     }
050     else {
051         hitColor = options.backgroundColor;
052     }
053
054     return hitColor;
055 }

```

Here is a result of rendering the scene with two spherical lights:



© www.scratchapixel.com

Note that you can get the same effect by rendering the contribution of each light individually and adding these images up in a paint program such as Photoshop (switch the blend mode to Linear Dodge). Check the render with the two lights on with your composite image in Photoshop. They should look the same.